

## Problem Statement :-

Build a model to accurately predict whether the patients in the dataset have diabetes or not?

## Data Description :-

The datasets consists of several medical predictor variables and one target variable, Outcome.

Predictor variables includes the number of pregnancies the patient has had, their BMI, insulin level, age, and so on.

Pregnancies: Number of times pregnant

Glucose: Plasma glucose concentration a 2 hours in an oral glucose tolerance test

BloodPressure: Diastolic blood pressure (mm Hg)

SkinThickness: Triceps skin fold thickness (mm)

Insulin: 2-Hour serum insulin (mu U/ml)

BMI: Body mass index (weight in kg/(height in m)^2)

DiabetesPedigreeFunction: Diabetes pedigree function

Age: Age (years)

Outcome: Class variable (0 or 1) 268 of 768 are 1, the others are 0

```
In [170...] import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

Importing the data

```
In [171...] data=pd.read_csv('/content/health care diabetes.csv')
```

```
In [172...] data.shape
```

```
Out[172...] (768, 9)
```

```
In [173...] data.columns
```

```
Out[173...] Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
      'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

```
In [174...] data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null    int64
1   Glucose                768 non-null    int64
2   BloodPressure          768 non-null    int64
3   SkinThickness          768 non-null    int64
4   Insulin               768 non-null    int64
5   BMI                   768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                   768 non-null    int64
8   Outcome               768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [175...] data['Outcome'].unique()
```

```
Out[175...] array([1, 0])
```

```
In [176...] # Last column of dataset i.e., Outcome is categorical because it has only two  
# unique items that are 0(not have diabetes) and 1(have diabetes) so I am changing int64 to category
```

```
In [177...] data['Outcome']=data['Outcome'].astype('category')
```

```
In [178...] data.info() # now Outcome is changes to category
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 768 entries, 0 to 767  
Data columns (total 9 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   Pregnancies            768 non-null   int64  
1   Glucose                768 non-null   int64  
2   BloodPressure          768 non-null   int64  
3   SkinThickness          768 non-null   int64  
4   Insulin                768 non-null   int64  
5   BMI                    768 non-null   float64  
6   DiabetesPedigreeFunction 768 non-null   float64  
7   Age                   768 non-null   int64  
8   Outcome                768 non-null   category  
dtypes: category(1), float64(2), int64(6)  
memory usage: 49.0 KB
```

## Checking null values in dataset

```
In [179...] data.isnull().sum()
```

```
Out[179...] Pregnancies            0  
Glucose                0  
BloodPressure          0  
SkinThickness          0  
Insulin                0  
BMI                    0  
DiabetesPedigreeFunction 0  
Age                   0  
Outcome                0  
dtype: int64
```

```
In [180...] # By this we clearly known that there are no null values in our dataset  
# But in some columns we have Zero values means missing values  
# Why because see in Glucose,BloodPressure,SkinThickness,Insulin,BMI zero means  
# As per domain knowledge Zero not makes sense so these Zero's are missing values
```

```
In [181...] lst = ['Glucose','BloodPressure','SkinThickness','Insulin','BMI']
```

```
In [182...] for i in lst:  
            print(i, (data[i]==0).sum())
```

```
Glucose 5  
BloodPressure 35  
SkinThickness 227  
Insulin 374  
BMI 11
```


```
In [183...] # By this we can see that missing values in Glucose,BloodPressure and BMI are negligible  
# But in SkinThickness and Insulin we have more missing values  
# I am not removing these missing values because our dataset is small dataset if we remove these  
# data points data becomes not sufficient so I am filing Zero's with Mean or Median
```

```
In [184...] # And Glucose=5, BloodPressure=35 and BMI=11 these have small missing values  
# so I am removing these rows with missing values
```

```
In [185...] data2 = data[(data["Glucose"] >0) & (data["BloodPressure"] >0) & (data['BMI']>0)]
```


```
In [186...] # First see outliers in these (SkinThickness and Insulin) two columns and decide which is best to fill
```

```
In [187...] plt.figure(figsize=(9,6))  
sns.boxplot(data2['SkinThickness'])  
plt.show()
```



It has less outliers so fill Zero's with mean

```
In [188... plt.figure(figsize=(9,6))
sns.boxplot(data2['Insulin'])
plt.show()
```



It has more outliers so fill Zero's with Median Because mean is affected by outliers and Median is not affected by outliers

```
In [189... data2['SkinThickness'].mean(), data2['Insulin'].median())
```

```
Out[189... (21.443370165745858, 48.0)
```


```
In [190... data2['SkinThickness']=data2['SkinThickness'].replace(0,21.443370165745858)
data2['Insulin']=data2['Insulin'].replace(0,48.0)
```

Detecting and Treating of outliers in dataset

```
In [191... data2.boxplot(figsize=(10, 10))
```

```
Out[191... <matplotlib.axes._subplots.AxesSubplot at 0x7fe47234ab50>
```






```
In [192]_ # By above boxplot Insulin have more outliers so we have to treat these outliers
# Removing outliers is not an option to our dataset because our dataset is small in size
```

```
In [193]_ figure=data2['Insulin'].hist(bins=50)
figure.set_title('Insulin')
figure.set_xlabel('Insulin')
figure.set_ylabel('No of Patients')
```

```
Out[193]_ Text(0, 0.5, 'No of Patients')
```



```
In [194]_ # Here the Data is not Normally distributed and is following the Right Skewed distribution
# that means we can use the Interquartile Range to measure the boundaries for outliers
# IQR = Q3 - Q1 (quartile 3 - quartile 1)
```

```
In [195]_ IQR= data2['Insulin'].quantile(0.75) - data2['Insulin'].quantile(0.25)
IQR
```

```
Out[195]_ 82.5
```

```
In [196]_ ## Calculating the boundaries
lower_bridge= data2['Insulin'].quantile(0.25)-(IQR*1.5)
upper_bridge= data2['Insulin'].quantile(0.75)+(IQR*1.5)
print(lower_bridge), print(upper_bridge)
```

```
-75.75
254.25
```

```
Out[196]_ (None, None)
```

```
In [197]_ data2['Insulin'].describe()
```

```
Out[197...] count    724.000000
mean     106.505525
std      102.669035
min       14.000000
25%       48.000000
50%       48.000000
75%      130.500000
max      846.000000
Name: Insulin, dtype: float64
```

```
In [198...] # Here the maximum value of outliers is very high compare to upper boundary that indicates
# we need to calculate the extreme outliers boundaries
```

```
In [199...] ## Calculating the extreme boundaries
lower_bridge= data2['Insulin'].quantile(0.25)-(IQR*3)
upper_bridge= data2['Insulin'].quantile(0.75)+(IQR*3)
print(lower_bridge), print(upper_bridge)

-199.5
378.0
```


```
Out[199...] (None, None)
```

```
In [200...] # Replacing the outliers
# Since the lower boundaries of Insulin column is negative value
# we do not need to consider the lower boundary because as per the domain knowledge
# there won't be any negative values exists for Insulin Column
# Insulin upper bridge is 378
```

```
In [201...] data2.loc[data2['Insulin']>=378,'Insulin']=378
```

```
In [202...] data2.boxplot(figsize=(10, 10))
```

```
Out[202...] <matplotlib.axes._subplots.AxesSubplot at 0x7fe472124710>
```




```
In [203...] # By comparing above boxplot outliers are only reduced not completely gone
```

Create a count (frequency) plot describing the data types and the count of variables.

```
In [204...] data2.dtypes.value_counts()
```

```
Out[204... int64      5
float64     3
category     1
dtype: int64
```

```
In [205... sns.countplot(data2.dtypes.map(str))
plt.show()
```




Check the balance of the data by plotting the count of outcomes by their value. Describe your findings and plan future course of actions.

```
In [206... data2['Outcome'].value_counts()
```

```
Out[206... 0      475
1      249
Name: Outcome, dtype: int64
```

```
In [207... sns.countplot(data2['Outcome'].map(str))
plt.show()
```




```
In [208... # By above we have clear evidence that our dataset has class imbalance issue
# There are different methods to solve class imbalance issue
# 1.Over Sampling Technique
# 2.Under sampling Technique
# 3.Cross validation Technique
# But see in both under and over samplung we are adding and deleting data points.
# High Accuracy will come but we cannot assure that that accuracy is biased or not
# But with cross validation in this techinque we train our model using the subset of the data-set
# so we can assure that accuracy got by this technique is unbaised.
```

Checking variables normal or not if not normal making to normal

```
In [209... sns.distplot(data2['Pregnancies'])
```

```
Out[209... <matplotlib.axes._subplots.AxesSubplot at 0x7fe471f03f50>
```






It is not normal it is right skewed

```
In [210...] sns.distplot(data2['Glucose'])
```

```
Out[210...] <matplotlib.axes._subplots.AxesSubplot at 0x7fe471ee2dd0>
```



it is visually normal but slight skewed so we can declare it is skewed or not by summary statistics


```
In [211...] data2['Glucose'].describe()
```

```
Out[211...] count    724.000000
mean      121.882597
std       30.750030
min       44.000000
25%      99.750000
50%     117.000000
75%     142.000000
max      199.000000
Name: Glucose, dtype: float64
```

Here there is 4 difference between mean and median so it is not normal

```
In [212...] sns.distplot(data2['BloodPressure'])
```

```
Out[212...] <matplotlib.axes._subplots.AxesSubplot at 0x7fe471dfadd0>
```



```
In [213...] data2['BloodPressure'].describe()
```


```
Out[213...] count    724.000000
mean       72.400552
std       12.379870
```

```
min      24.000000
25%      64.000000
50%      72.000000
75%      80.000000
max     122.000000
Name: BloodPressure, dtype: float64
```

It is normal

```
In [214...] sns.distplot(data2['SkinThickness'])
```

```
Out[214...] <matplotlib.axes._subplots.AxesSubplot at 0x7fe471d08e10>
```




```
In [215...] data2['SkinThickness'].describe()
```

```
Out[215...] count      724.000000
mean       27.130010
std        9.645083
min         7.000000
25%        21.443370
50%        24.000000
75%        33.000000
max        99.000000
Name: SkinThickness, dtype: float64
```

Not normal right skewed

```
In [216...] sns.distplot(data2['Insulin'])
```

```
Out[216...] <matplotlib.axes._subplots.AxesSubplot at 0x7fe471c43090>
```



```
In [217...] data2['Insulin'].describe()
```

```
Out[217...] count      724.000000
mean       102.142265
std        84.536572
min        14.000000
25%        48.000000
50%        48.000000
75%       130.500000
```




```
max      378.000000
Name: Insulin, dtype: float64
```

Not normal and right skewed

```
In [218...] sns.distplot(data2['BMI'])
```

```
Out[218...] <matplotlib.axes._subplots.AxesSubplot at 0x7fe471b75190>
```



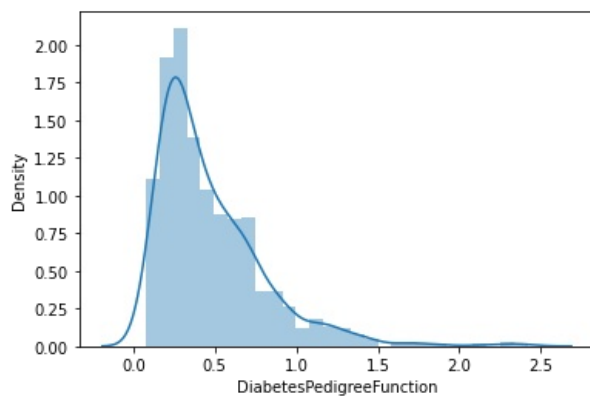
```
In [219...] data2['BMI'].describe()
```

```
Out[219...] count      724.000000
mean        32.467127
std         6.888941
min         18.200000
25%         27.500000
50%         32.400000
75%         36.600000
max         67.100000
Name: BMI, dtype: float64
```

It is normal distribution

```
In [220...] sns.distplot(data2['DiabetesPedigreeFunction'])
```

```
Out[220...] <matplotlib.axes._subplots.AxesSubplot at 0x7fe471b10350>
```




```
In [221...] data2['DiabetesPedigreeFunction'].describe()
```

```
Out[221...] count      724.000000
mean         0.474765
std          0.332315
min          0.078000
25%          0.245000
50%          0.379000
75%          0.627500
max          2.420000
Name: DiabetesPedigreeFunction, dtype: float64
```

It is normal difference between mean and median is negligible

```
In [222...] sns.distplot(data2['Age'])
```

```
Out[222...] <matplotlib.axes._subplots.AxesSubplot at 0x7fe471a38990>
```



```
In [223...] data2['Age'].describe()
```

```
Out[223...] count    724.000000
mean      33.350829
std       11.765393
min       21.000000
25%       24.000000
50%       29.000000
75%       41.000000
max       81.000000
Name: Age, dtype: float64
```

It is not normal

So here Pregnancies,Glucose,SkinThickness,Insulin,Age are not normal we have to make it to normal

```
In [224...] lst1 = ['Pregnancies','Glucose','SkinThickness','Insulin','Age']
```

```
In [225...] for i in lst1:
data2[i] = np.log1p(data2[i])
```

```
In [226...] data2.describe()
```

```
Out[226...]      Pregnancies    Glucose  BloodPressure  SkinThickness    Insulin      BMI  DiabetesPedigreeFunction    Age
count    724.000000    724.000000    724.000000    724.000000    724.000000    724.000000    724.000000    724.000000
mean      1.319311    4.780277      72.400552      3.280856    4.387084    32.467127      0.474765    3.484730
std       0.762929    0.249781     12.379870     0.336677    0.662871     6.888941     0.332315    0.313971
min       0.000000    3.806662     24.000000     2.079442    2.708050    18.200000     0.078000    3.091042
25%       0.693147    4.612633     64.000000     3.110995    3.891820    27.500000     0.245000    3.218876
50%       1.386294    4.770685     72.000000     3.218876    3.891820    32.400000     0.379000    3.401197
75%       1.945910    4.962845     80.000000     3.526361    4.878985    36.600000     0.627500    3.737670
max       2.890372    5.298317    122.000000     4.605170    5.937536    67.100000     2.420000    4.406719
```

Checking correlation between variables

First Doing Scaling to all dataset

```
In [227...] from sklearn.preprocessing import MinMaxScaler,StandardScaler
m=MinMaxScaler()
```

Here x is independent variable and y is dependent variable (Outcome)

```
In [228...] x=data2.drop(['Outcome'],axis=1)
y=data2['Outcome']
```

```
In [229] data_sc=m.fit_transform(x)
```

```
In [230] data_sc_df=pd.DataFrame(data_sc,columns=x.columns,index=x.index)
data_sc_df.head(5)
```

Out[230]

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	0.673239	0.802655	0.489796	0.595502	0.366551	0.314928	0.234415	0.639050
1	0.239812	0.434206	0.428571	0.523317	0.366551	0.171779	0.116567	0.284791
2	0.760188	0.944101	0.408163	0.408418	0.366551	0.104294	0.253629	0.308180
3	0.239812	0.464683	0.428571	0.434968	0.571554	0.202454	0.038002	0.000000
4	0.000000	0.751240	0.163265	0.595502	0.749918	0.509202	0.943638	0.330870

Finding correlation between variables


```
In [231] data_sc_df.corr()
```

Out[231]

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
Pregnancies	1.000000	0.110236	0.159556	0.000653	-0.061777	-0.046226	-0.046698	0.560507
Glucose	0.110236	1.000000	0.226538	0.164075	0.354076	0.220961	0.125562	0.257609
BloodPressure	0.159556	0.226538	1.000000	0.147366	-0.035981	0.287403	-0.000075	0.340852
SkinThickness	0.000653	0.164075	0.147366	1.000000	0.249853	0.561603	0.119397	0.069261
Insulin	-0.061777	0.354076	-0.035981	0.249853	1.000000	0.205610	0.167646	-0.035048
BMI	-0.046226	0.220961	0.287403	0.561603	0.205610	1.000000	0.154858	0.049484
DiabetesPedigreeFunction	-0.046698	0.125562	-0.000075	0.119397	0.167646	0.154858	1.000000	0.032301
Age	0.560507	0.257609	0.340852	0.069261	-0.035048	0.049484	0.032301	1.000000

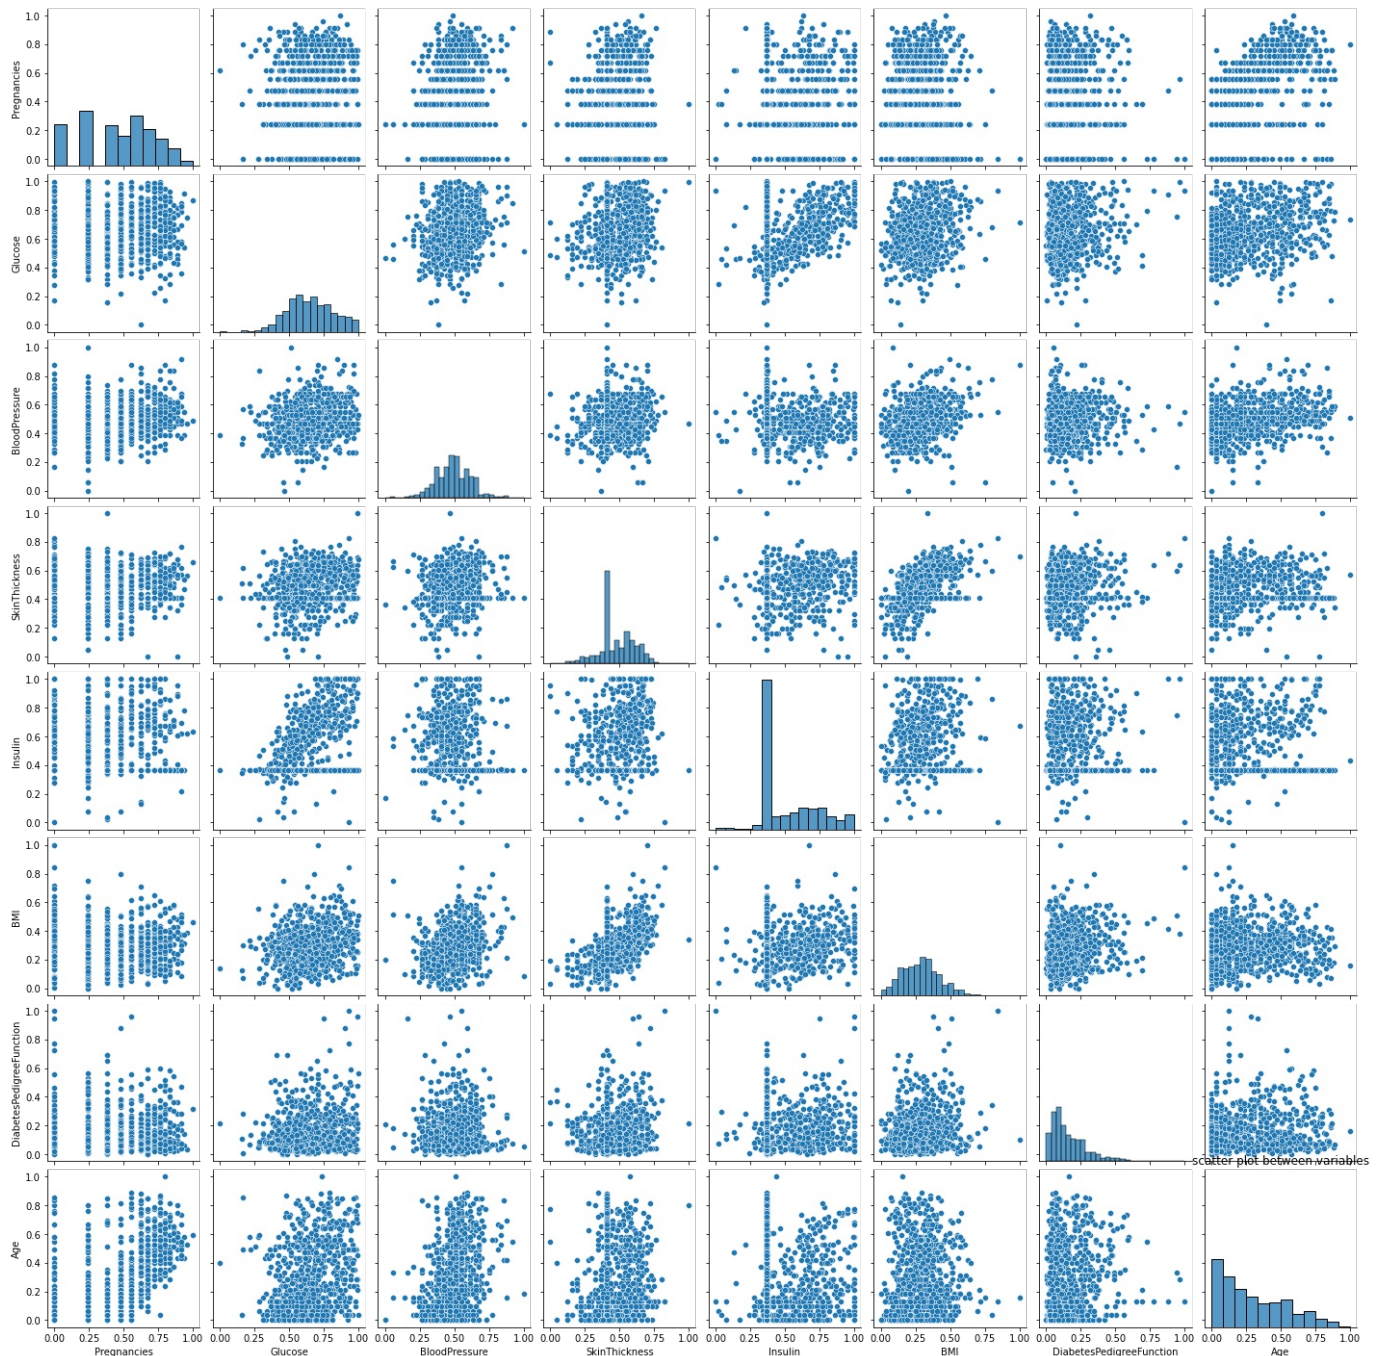
```
In [232] plt.figure(figsize = (9,9))
sns.heatmap(data_sc_df.corr(),annot=True,fmt='.2f')
```

Out[232] <matplotlib.axes.\_subplots.AxesSubplot at 0x7fe4719c6f10>



```
In [233]: sns.pairplot(data_sc_df)
plt.title('scatter plot between variables')
```

```
Out[233]: Text(0.5, 1.0, 'scatter plot between variables')
```



```
In [234]: # Clearly multi-collinearity exists in variables in our data set
# We can see from scatter plot that there is no strong multicollinearity among features
# But between skin thickness and BMI, Pregnancies and age it looks like there positive correlation.
# In heat map we got cleared there is multi-collinearity
```

```
In [235]: # For Logistic Regression multi-collinearity should not be in data
# But here there is clearly multi-collinearity exists between variables
# Multi-collinearity means presence of relation between independent variables
# Multi-collinearity is a problem we should avoid
# There are two types o avoid multi-collinearity
# 1.) Remove independent variables which are corelated
# 2.) Use dimensionality reduction technique to merge variables
```

```
In [236]: # Variance Inflation Factor (VIF)
# equal to the ratio of the overall model variance to the variance of a model
# that includes only that single independent variable.
# VIF=1/1-R square (R-square is correlation coefficient of that variable)
```

```
In [237]: # VIF ~ 1: Negligible
# 1<VIF<5 : Moderate
```

```
# VIF>5 : Extreme
# if a variable with extreme VIF is should be removed
```

```
In [238...] from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [239...] # VIF dataframe
vif_data = pd.DataFrame()
vif_data["feature"] = data_sc_df.columns
```

```
In [240...] # calculating VIF for each feature
vif_data["VIF"] = [variance_inflation_factor(data_sc_df.values, i)
for i in range(len(data_sc_df.columns))]
```

```
In [241...] print(vif_data)
```

	feature	VIF
0	Pregnancies	5.512882
1	Glucose	18.220714
2	BloodPressure	15.132647
3	SkinThickness	17.090657
4	Insulin	9.077077
5	BMI	8.266541
6	DiabetesPedigreeFunction	2.534800
7	Age	4.104236

```
In [242...] # Here five variables have high VIF value so we have to remove these variables
# That leads heavy data loss in our scenario so by using PCA technique.
# we are going to merge there all variables.
```

PCA Technique

```
In [243...] # using PCA technique
from sklearn.decomposition import PCA
```

```
In [244...] pca=PCA(n_components=6)
var_x=pca.fit_transform(data_sc_df)
```

```
In [245...] new_data=pd.DataFrame(var_x,columns=['pca1','pca2','pca3','pca4','pca5','pca6'])
```


```
In [246...] b=pca.explained_variance_ratio_
b*100
```

```
Out[246...] array([37.81312702, 20.79637505, 11.53703291,  8.99408893,  7.07498587,
 6.66814181])
```

Here 88% data captured by these pca variables

```
In [247...] plt.figure(figsize = (5,5))
sns.heatmap(new_data.corr())
```

```
Out[247...] <matplotlib.axes._subplots.AxesSubplot at 0x7fe46fdc9a90>
```



so finally no multicollinearity so we can proceed with logistic regression

## Logistic Regression Model Building

## Logistic Regression Model Building

```
In [248... # so now
x_var=new_data
y_var=y
```

```
In [249... # splitting our data
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score,classification_report
```

```
In [250... x_train,x_test,y_train,y_test=train_test_split(x_var,y_var,test_size=0.2,random_state=50)
```

```
In [251... from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()
lr.fit(x_train,y_train)
```

```
Out[251... LogisticRegression()
```

```
In [252... pred=lr.predict(x_test)
```

```
In [253... accuracy_score(y_test,pred)
```

```
Out[253... 0.7793103448275862
```

```
In [254... accuracy_score(y_train,lr.predict(x_train))
```

```
Out[254... 0.7651122625215889
```

```
In [255... from prettytable import PrettyTable
x=PrettyTable(['Accuracy','score'])
x.add_row(['y_test and y_test predicted',0.7793103448275862])
x.add_row(['y_train and y_train predicted',0.7651122625215889])
print(x)
```

```
+-----+-----+
|          Accuracy          |          score          |
+-----+-----+
| y_test and y_test predicted | 0.7793103448275862 |
| y_train and y_train predicted | 0.7651122625215889 |
+-----+-----+
```

see clearly our model is not over fitted so see for other reports

```
In [256... print(classification_report(y_test,pred))
```

	precision	recall	f1-score	support
0	0.76	0.94	0.84	89
1	0.85	0.52	0.64	56
accuracy			0.78	145
macro avg	0.80	0.73	0.74	145
weighted avg	0.79	0.78	0.76	145

1.Accuracy is over all 78% percent it is good because our data predicts 78% correctly 2.But here our outcome is having diabetes (1) and not having diabetes (0) means both 1 and 0 are important here so we have to consider macro average because it gives equal importance to both 1 and 0 so our macro average is 74%

## AUC and ROC curve

### AUC Curve

The Area Under the Curve (AUC) Is the measure of the ability of a classifier to distinguish between classes and is used as a summary of the ROC curve. The higher the AUC, the better the performance of the model at distinguishing between the positive and negative class

### ROC curve

ROC (Receiver Operating Characteristic Curve) In ROC curve we plot pairs of the true positive (Sensitivity) rate vs the false positive rate (Specificity), for every possible decision threshold of a logistic regression model. By ROC curve we can find optimal Threshold value for our model.

```
In [257... # plot roc and auc curve
from sklearn.metrics import roc_auc_score, roc_curve
```

Here first find predicted probabilities of test data. Then created dataframe of actual and predicted probabilities of data


```
In [258... y_prob=lr.predict_proba(x_test)
df_pred_prob=pd.DataFrame({'Actual':y_test,'Predicted_prob':y_prob[:,1]})
```

```
In [259... df_pred_prob
```

```
Out[259...
   Actual Predicted_prob
19      1         0.274795
750     1         0.396869
213     1         0.318996
424     1         0.713071
501     0         0.144898
...     ...           ...
647     1         0.405455
108     0         0.088859
730     1         0.315705
761     1         0.726563
45      1         0.704235
```

145 rows × 2 columns

```
In [260... y_prob[:,1][0:10]
fpr, tpr, thresholds=roc_curve(y_test,y_prob[:,1])
auc_curve=roc_auc_score(y_test,y_prob[:,1])
plt.plot(fpr,tpr,label="AUC="+str(auc_curve))
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.legend(loc=4)
plt.show()
```



Here our AUC is 85% means our model distinguished 85% correctly

```
In [261... from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
confusion_matrix(y_test,pred)
```

```
Out[261... array([[84,  5],
       [27, 29]])
```

BY this confusion matrix we can say that actually 5 patients have no diabetes but our model says have diabetes. Actually 27 patients have diabetes but our predicted model says not have diabetes this is false negative here our false negative=27 we have to reduce our False Negative. we have to find threshold value to reduce our error(false negative) by ROC curve we can know that

But threshold at 0.4 means probabilities more than 0.4 are 1 (have diabetes) rest are 0 (not have diabetes)



Put threshold at 0.4 means probabilities more than 0.4 are 1 (have diabetes) rest are 0 (not have diabetes)

```
In [262... df_pred_prob['prediction_0.4']=np.where(df_pred_prob['Predicted_prob']>0.4,1,0)
confusion_matrix(df_pred_prob['Actual'],df_pred_prob['prediction_0.4'])
```

```
Out[262... array([[78, 11],
        [18, 38]])
```

Here our false negative is reduced I think it's fair so I am fixing threshold value at 0.4 . By domain knowledge if we want min (least) false negative we can reduce threshold value below 0.4

put threshold at 0.2

```
In [263... df_pred_prob['prediction_0.2']=np.where(df_pred_prob['Predicted_prob']>0.2,1,0)
confusion_matrix(df_pred_prob['Actual'],df_pred_prob['prediction_0.2'])
```

```
Out[263... array([[46, 43],
        [ 4, 52]])
```

Here see false negative(FN) is 4 means FN is reduced but see False positive(FP) is increased

## Other Classification Models Except Logistic Model

### Decision Trees

For Tree models preprocessing steps are not necessary but we will use above data

```
In [264... # Here I am using Scaling Data not PCA data
x_dtree=data_sc_df
y_dtree=y
```

```
In [265... x_dtree_train,x_dtree_test,y_dtree_train,y_dtree_test=train_test_split(x_dtree,y_dtree,test_size=0.2,random_state=50)
```

```
In [266... from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier()
dt.fit(x_dtree_train,y_dtree_train)
dpred=dt.predict(x_dtree_test)
```

```
In [267... from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
```

```
In [268... accuracy_score(y_dtree_test,dpred)
```

```
Out[268... 0.6689655172413793
```

```
In [269... accuracy_score(y_dtree_train,dt.predict(x_dtree_train))
```

```
Out[269... 1.0
```

Now use PCA independent variables and check accuracy

```
In [270... x_train_p,x_test_p,y_train_p,y_test_p=train_test_split(x_var,y_var,test_size=0.2,random_state=50)
```

```
In [271... from sklearn.tree import DecisionTreeClassifier
pdt=DecisionTreeClassifier()
```

```
In [272... pdt.fit(x_train_p,y_train_p)
```

```
Out[272... DecisionTreeClassifier()
```

```
In [273... pdpred=pdt.predict(x_test_p)
```

```
In [274... accuracy_score(y_test_p,pdpred)
```

```
Out[274... 0.7034482758620689
```



```
In [275...] accuracy_score(y_train_p,pdt.predict(x_train_p))
```

```
Out[275...] 1.0
```

There is small difference in accuracy between scaling data and PCA data (more accuracy in PCA data). So we can use PCA independent variables for further models because our PCA data is multicollinearity free and it reduces overfitting. Here our model is clearly over fitted. It is I think due to class imbalance issue. First we build all classification models and finally we see every model with cross validation technique and finalize our model based on accuracy got by cross validation technique.

## Random Forest

```
In [276...] x_rforest=x_var  
y_rforest=y_var
```

```
In [277...] from sklearn.ensemble import RandomForestClassifier  
rf=RandomForestClassifier()
```

```
In [278...] x_rforest_train,x_rforest_test,y_rforest_train,y_rforest_test=train_test_split(x_rforest,y_rforest,test_size=0.2,
```

```
In [279...] rf.fit(x_rforest_train,y_rforest_train)
```

```
Out[279...] RandomForestClassifier()
```

```
In [280...] rf_pred=rf.predict(x_rforest_test)
```

```
In [281...] accuracy_score(y_rforest_test,rf_pred)
```

```
Out[281...] 0.7517241379310344
```

```
In [282...] accuracy_score(y_rforest_train,rf.predict(x_rforest_train))
```

```
Out[282...] 1.0
```

```
In [283...] print(classification_report(y_rforest_test,rf_pred))
```

	precision	recall	f1-score	support
0	0.74	0.91	0.82	89
1	0.78	0.50	0.61	56
accuracy			0.75	145
macro avg	0.76	0.71	0.71	145
weighted avg	0.76	0.75	0.74	145

## Support Vector Machine Model

```
In [284...] x_svm=x_var  
y_svm=y_var
```

```
In [285...] x_svm_train,x_svm_test,y_svm_train,y_svm_test=train_test_split(x_svm,y_svm,test_size=0.2,random_state=50)
```

```
In [286...] from sklearn.svm import SVC
```

```
In [287...] svm=SVC()  
svm.fit(x_svm_train,y_svm_train)
```

```
Out[287...] SVC()
```

```
In [288...] svm_pred=svm.predict(x_svm_test)
```

```
In [289...] accuracy_score(y_svm_test,svm_pred)
```

Out[289...] 0.7241379310344828

```
In [290...] accuracy_score(y_svm_train,svm.predict(x_svm_train))
```

Out[290...] 0.7962003454231433

```
In [291...] print(classification_report(y_svm_test,svm_pred))
```

	precision	recall	f1-score	support
0	0.72	0.91	0.80	89
1	0.75	0.43	0.55	56
accuracy			0.72	145
macro avg	0.73	0.67	0.67	145
weighted avg	0.73	0.72	0.70	145

## Naive Bayes Classification Algorithm

```
In [292...] x_train_nb,x_test_nb,y_train_nb,y_test_nb=train_test_split(data_sc_df,y_var,test_size=0.2,random_state=50)
```

```
In [293...] from sklearn.naive_bayes import MultinomialNB  
nb=MultinomialNB()
```

```
In [294...] nb.fit(x_train_nb,y_train_nb)
```

Out[294...] MultinomialNB()

```
In [295...] pred_nb=nb.predict(x_test_nb)
```

```
In [296...] accuracy_score(y_test_nb,pred_nb)
```

Out[296...] 0.6137931034482759

```
In [297...] accuracy_score(y_train_nb,nb.predict(x_train_nb))
```

Out[297...] 0.6666666666666666

```
In [298...] print(classification_report(y_test_nb,pred_nb))
```

	precision	recall	f1-score	support
0	0.61	1.00	0.76	89
1	0.00	0.00	0.00	56
accuracy			0.61	145
macro avg	0.31	0.50	0.38	145
weighted avg	0.38	0.61	0.47	145

## K-Nearest Neighbors Algorithm

I am using PCA independent variables

```
In [299...] x_var.head(2)
```

	pca1	pca2	pca3	pca4	pca5	pca6
0	0.407391	0.001247	0.171459	-0.012504	0.031812	0.080258
1	-0.205940	0.217559	0.105583	0.029990	0.076450	-0.148843

```
In [300... from sklearn.cluster import KMeans
km = KMeans(n_clusters=5)
km.fit(x_var)
```

```
Out[300... KMeans(n_clusters=5)
```

```
In [301... km.labels_
```

```
Out[301... array([1, 2, 1, 2, 4, 2, 2, 3, 1, 1, 1, 3, 3, 4, 1, 0, 0, 4, 1, 1, 1, 3,
      3, 1, 2, 1, 1, 1, 4, 2, 2, 1, 3, 1, 1, 2, 3, 2, 1, 1, 3, 1, 0, 2,
      2, 1, 2, 2, 2, 3, 3, 2, 3, 0, 0, 0, 1, 2, 4, 1, 1, 0, 1, 2, 2, 2,
      4, 1, 4, 2, 1, 2, 2, 2, 1, 0, 1, 4, 1, 2, 3, 2, 2, 3, 1, 1, 2, 3,
      2, 2, 2, 4, 0, 2, 0, 2, 2, 4, 0, 3, 2, 0, 4, 3, 2, 2, 3, 1, 1, 2,
      2, 2, 0, 2, 2, 1, 0, 4, 4, 4, 4, 0, 3, 1, 4, 1, 2, 4, 0, 0, 0, 4,
      1, 1, 2, 1, 4, 1, 4, 1, 2, 4, 1, 3, 4, 1, 1, 2, 4, 2, 3, 1, 1, 4,
      2, 0, 3, 2, 1, 2, 2, 1, 3, 0, 2, 3, 1, 0, 1, 1, 1, 0, 2, 1, 1, 3,
      0, 3, 4, 2, 1, 1, 1, 3, 2, 2, 4, 4, 0, 0, 0, 2, 3, 2, 3, 1, 2, 1,
      2, 0, 1, 0, 3, 3, 4, 3, 2, 1, 4, 1, 3, 2, 2, 0, 2, 3, 0, 2, 3, 2,
      2, 2, 2, 3, 0, 1, 0, 2, 2, 2, 3, 4, 1, 1, 4, 3, 2, 1, 2, 2, 0, 3,
      0, 2, 2, 4, 3, 3, 2, 1, 1, 1, 0, 0, 1, 2, 1, 0, 1, 0, 2, 0, 1, 4,
      0, 3, 3, 1, 1, 3, 3, 4, 2, 1, 0, 0, 4, 4, 0, 3, 4, 4, 3, 1, 4, 1,
      1, 1, 4, 3, 0, 4, 4, 1, 0, 2, 2, 1, 2, 2, 2, 4, 1, 3, 2, 0, 1, 2,
      4, 4, 1, 4, 1, 1, 2, 1, 2, 4, 1, 3, 1, 2, 2, 1, 1, 3, 2, 2, 2, 1,
      1, 2, 2, 1, 4, 1, 4, 3, 1, 1, 1, 3, 2, 1, 0, 2, 4, 4, 0, 2, 4, 3,
      0, 0, 1, 0, 0, 0, 2, 2, 2, 2, 1, 1, 3, 2, 4, 1, 4, 1, 1, 4, 3, 0,
      2, 2, 2, 1, 1, 1, 1, 4, 1, 0, 1, 4, 1, 4, 4, 2, 4, 4, 2, 1, 2, 2,
      4, 2, 0, 2, 3, 3, 4, 0, 4, 2, 2, 2, 2, 1, 2, 2, 1, 0, 2, 2, 1, 2,
      0, 2, 0, 0, 0, 2, 2, 4, 2, 1, 1, 2, 3, 1, 1, 2, 1, 1, 0, 2, 0,
      3, 0, 0, 0, 1, 2, 0, 4, 1, 1, 1, 3, 0, 2, 0, 4, 4, 4, 2, 1, 2, 1,
      1, 3, 1, 2, 2, 3, 3, 2, 2, 1, 1, 1, 0, 4, 2, 1, 1, 4, 1, 2, 2, 4,
      3, 1, 1, 3, 2, 4, 1, 2, 2, 2, 2, 4, 0, 2, 0, 0, 2, 0, 0, 4, 4, 3,
      4, 1, 2, 0, 3, 3, 4, 1, 1, 2, 2, 1, 2, 0, 3, 0, 1, 1, 1, 1, 4, 0,
      2, 0, 2, 0, 1, 3, 4, 1, 2, 2, 2, 4, 0, 3, 2, 1, 1, 0, 2, 1, 1, 3,
      2, 1, 2, 3, 1, 4, 1, 2, 3, 4, 0, 2, 0, 2, 2, 0, 3, 0, 4, 2, 4, 2,
      2, 3, 3, 2, 3, 2, 1, 2, 1, 4, 2, 1, 0, 2, 2, 0, 0, 1, 2, 1, 0, 2,
      4, 1, 1, 1, 2, 1, 2, 0, 2, 1, 4, 4, 4, 4, 3, 0, 2, 4, 2, 2, 4, 4,
      2, 4, 1, 2, 1, 0, 3, 3, 1, 4, 1, 1, 3, 3, 3, 2, 1, 4, 1, 1, 1, 0,
      2, 4, 2, 0, 0, 2, 4, 2, 2, 4, 4, 1, 1, 4, 3, 2, 3, 4, 4, 2, 4, 1,
      1, 2, 2, 4, 1, 4, 4, 1, 1, 4, 2, 3, 4, 1, 4, 1, 1, 4, 4, 1, 1, 1,
      4, 0, 2, 2, 1, 2, 4, 4, 1, 2, 0, 1, 4, 0, 3, 2, 2, 1, 3, 1, 0, 0,
      3, 1, 2, 0, 2, 4, 1, 4, 1, 0, 0, 1, 2, 1, 1, 3, 2, 2, 1, 2],
      dtype=int32)
```

```
In [302... d=x_var
d["labels"] =km.labels_
```

```
In [303... ## elbow technique to decide the number of clusters
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [304... wcss=[]
for i in range(1,14):
    km = KMeans(n_clusters=i, init="k-means++")
    km.fit(x_var)
    wcss.append(km.inertia_)
plt.figure(figsize=(12,6))
sns.lineplot(range(1,14), wcss, color="red")
```

```
Out[304... <matplotlib.axes._subplots.AxesSubplot at 0x7fe4721c29d0>
```





```
In [305... km.inertia_
Out[305... 62.401864862006605
```

See by inertia at 5 clusters we got low inertia that is 62% so our data distinguished into 5 clusters

## Model Selection

Using Cross Validation Technique to finalize model

- 1.I am using cross validation technique other than over and under sampling.
- 2.Because In over and under sampling there is a chance that accuracy is biased because we are adding data points in over sampling and removing some data points in under sampling.
- 3.But in cross validation technique this algorithm takes random sample subsets in dataset as test data and goes through further analysis so in this process there is less chance of overfitting and we can get unbiased accuracy.
- 4.And in this technique it calculates accuracies of all subsets and takes its average so means it considers all values in data even data is in imbalance.

### Logistic Regression Cross Validation

Cross validation technique for logistic regression where we used PCA variables as independent variables.

```
In [306... from sklearn.model_selection import cross_validate
I am taking 5 subsets

In [307... res_lr=cross_validate(lr,x_var,y_var,cv=5,return_train_score=True)
Test Score

In [308... test_lr=np.average(res_lr['test_score'])
test_lr
Out[308... 0.7638314176245211
```

```
In [309... train_lr=np.average(res_lr['train_score'])
train_lr
Out[309... 0.7710612828300876
```

### Decision Tree Cross Validation

I am taking 5 subsets

```
In [310... res_dt=cross_validate(pdt,x_var,y_var,cv=5,return_train_score=True)
Test Score

In [311... test_dt=np.average(res_dt['test_score'])
test_dt
Out[311... 0.6615229885057472
```

```
In [312... train_dt=np.average(res_dt['train_score'])
train_dt
```

Out[312...] 1.0

## Random Forest Cross Validation

I am taking 4 subsets

```
In [313...] res_rf=cross_validate(rf,x_rforest,y_rforest,cv=4,return_train_score=True)
```

Test Score

```
In [314...] test_rf=np.average(res_rf['test_score'])
test_rf
```

Out[314...] 0.7527624309392265

```
In [315...] train_rf=np.average(res_rf['train_score'])
train_rf
```

Out[315...] 1.0

## Support Vector Machine Model Cross Validation

I am taking 5 subsets

```
In [316...] res_svm=cross_validate(svm,x_svm,y_svm,cv=5,return_train_score=True)
```

Test Score

```
In [317...] test_svm=np.average(res_svm['test_score'])
test_svm
```

Out[317...] 0.7389463601532567

```
In [318...] train_svm=np.average(res_svm['train_score'])
train_svm
```

Out[318...] 0.7569060806384373

## Naive Bayes Classification Algorithm Model Cross Validation

I am taking 5 subsets

```
In [319...] res_nb=cross_validate(nb,data_sc_df,y_var,cv=5,return_train_score=True)
```

Test Score

```
In [320...] test_nb=np.average(res_nb['test_score'])
test_nb
```

Out[320...] 0.6560823754789272

Train Score

```
In [321...] train_nb=np.average(res_nb['train_score'])
train_nb
```

Out[321...] 0.6560776606515395

## Comparing Models Based On Cross Validation Technique

---

```
In [322... from prettytable import PrettyTable
x=PrettyTable(['Model','Train_score','Test_score'])
x.add_row(['Logistic Regression',0.7665713775236733,0.7569061302681993])
x.add_row(['Decision Tree',1.0,0.6587164750957853])
x.add_row(['Random Forest',1.0,0.7596685082872928])
x.add_row(['Support Vector Machine',0.7272098147817283,0.7126819923371647])
x.add_row(['Naive Bayes Classification Algorithm',0.6560776606515395,0.6560823754789272])
print(x)
```

Model	Train_score	Test_score
Logistic Regression	0.7665713775236733	0.7569061302681993
Decision Tree	1.0	0.6587164750957853
Random Forest	1.0	0.7596685082872928
Support Vector Machine	0.7272098147817283	0.7126819923371647
Naive Bayes Classification Algorithm	0.6560776606515395	0.6560823754789272

1. There is only 1% difference in accuracy of train and test data in Logistic Regression.
2. Among all models logistic regression got more accuracy and overfitting also less so we fix logistic regression as our final model.

## Finding Best Parameters For Logistic Regression By Hyperparametric Tuning Technique

```
In [323... # Gridsearchcv
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
```

```
In [324... pca=PCA(n_components=6)
x_vari=pca.fit_transform(data_sc_df)
```

```
In [328... new_data1=pd.DataFrame(var_x,columns=['pca1','pca2','pca3','pca4','pca5','pca6'])
```

```
In [332... x_traing,x_testg,y_traing,y_testg=train_test_split(new_data1,y_var,test_size=0.2,random_state=50)
```

```
In [341... grid = {'C' : np.logspace(-3,3,7), 'penalty' : ['l1','l2']}
logreg = LogisticRegression()
logreg_cv = GridSearchCV(logreg,grid,cv = 10)
logreg_cv.fit(x_traing,y_traing)
print("Tuned hyperparameters : (best parameters)", logreg_cv.best_params_)
print("accuracy : ",logreg_cv.best_score_)
```

Tuned hyperparameters : (best parameters) {'C': 10.0, 'penalty': 'l2'}

accuracy : 0.7686025408348457

By this also we proved that logistic regression is best because we got 76% accuracy in our model also we got 0.7569 ~ 76 so we predicted correctly and above parameters are best parameters

## Final Model With Function

```
In [ ]: from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score,classification_report
from sklearn.linear_model import LogisticRegression
```

Using Logistic Model

```
In [352... def predict_class_logistic(x_train,x_test,y_train,y_test):
    """this function predicts the final o/p class label"""
    """ put scaled data to form PCA after all necessary preprocessing steps and insert in model"""
    model = LogisticRegression(penalty = 'l2',C = 10.0)
    model.fit(x_train,y_train)
    pred = model.predict(x_test)
    return roc_auc_score(y_test,pred).round(2)
```

```
In [353... predict_class_logistic(x_traing,x_testg,y_traing,y_testg)
```

Out[353... 0.74

By this we defined our model function